

Diabetes-Risk-Prediction (/github/ShubhaTiwarii/Diabetes-Risk-Prediction/tree/main)

/

Diabetes Risk Analysis.ipynb (/github/ShubhaTiwarii/Diabetes-Risk-Prediction/tree/main/Diabetes Risk Analysis.ipynb)

In [2]: *#Diabetes Risk Analysis and Prediction based on different symptoms, age and gender*

```
import pandas as pd
import numpy as np
df = pd.read_csv("diabetes_data_upload.csv")
df.head(10)
```

Out[2]:

	Age	Gender	Polyuria	Polydipsia	sudden weight loss	weakness	Polyphagia	Genital thrush	visual blurring	Itching	Irritabili
0	40	Male	No	Yes	No	Yes	No	No	No	Yes	↑
1	58	Male	No	No	No	Yes	No	No	Yes	No	↑
2	41	Male	Yes	No	No	Yes	Yes	No	No	Yes	↑
3	45	Male	No	No	Yes	Yes	Yes	Yes	No	Yes	↑
4	60	Male	Yes	Yes	Yes	Yes	Yes	No	Yes	Yes	Y
5	55	Male	Yes	Yes	No	Yes	Yes	No	Yes	Yes	↑
6	57	Male	Yes	Yes	No	Yes	Yes	Yes	No	No	↑
7	66	Male	Yes	Yes	Yes	Yes	No	No	Yes	Yes	Y
8	67	Male	Yes	Yes	No	Yes	Yes	Yes	No	Yes	Y
9	70	Male	No	Yes	Yes	Yes	Yes	No	Yes	Yes	Y

In [3]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 520 entries, 0 to 519
Data columns (total 17 columns):
#   Column                      Non-Null Count  Dtype
---  -
0   Age                          520 non-null    int64
1   Gender                       520 non-null    object
2   Polyuria                     520 non-null    object
3   Polydipsia                   520 non-null    object
4   sudden weight loss           520 non-null    object
5   weakness                     520 non-null    object
6   Polyphagia                   520 non-null    object
7   Genital thrush               520 non-null    object
8   visual blurring              520 non-null    object
9   Itching                      520 non-null    object
10  Irritability                 520 non-null    object
11  delayed healing              520 non-null    object
12  partial paresis              520 non-null    object
13  muscle stiffness             520 non-null    object
14  Alopecia                     520 non-null    object
15  Obesity                      520 non-null    object
16  class                        520 non-null    object
dtypes: int64(1), object(16)
memory usage: 69.2+ KB
```

```
In [4]: df.isnull().sum()
```

```
Out[4]: Age                0
Gender                0
Polyuria              0
Polydipsia            0
sudden weight loss    0
weakness              0
Polyphagia            0
Genital thrush        0
visual blurring       0
Itching               0
Irritability          0
delayed healing        0
partial paresis        0
muscle stiffness       0
Alopecia              0
Obesity               0
class                 0
dtype: int64
```

```
In [5]: df.shape
```

```
Out[5]: (520, 17)
```

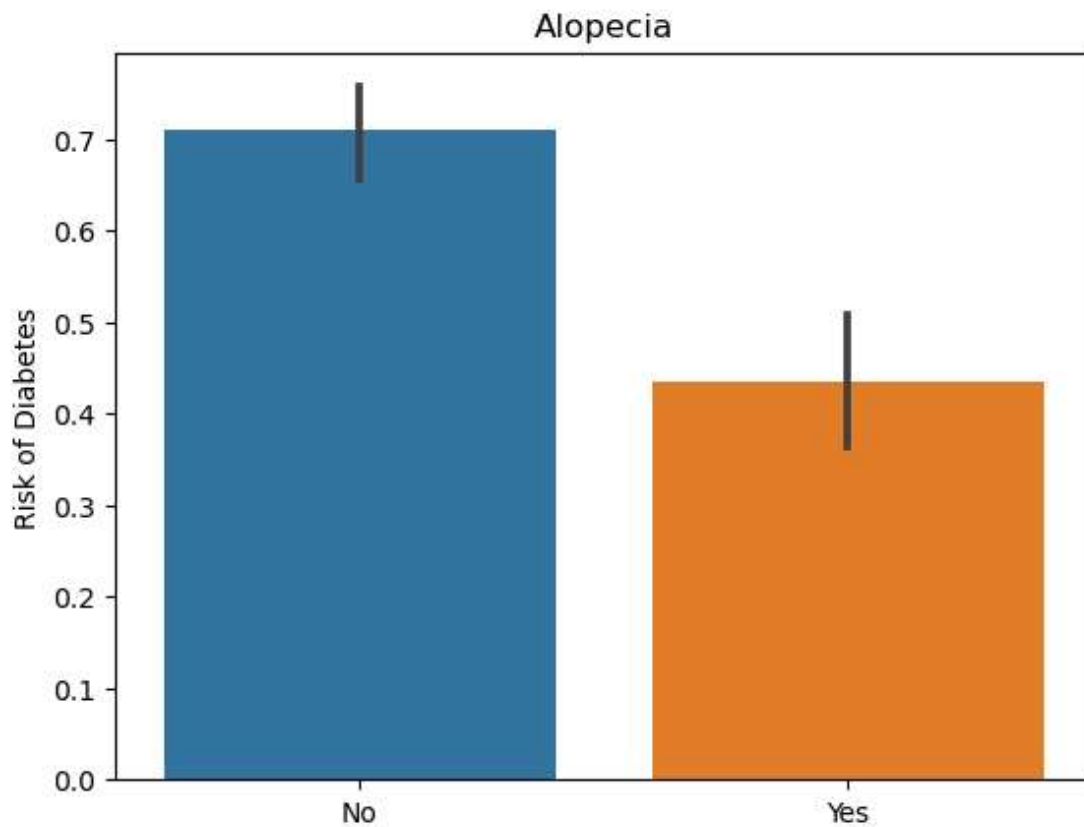
```
In [6]: df1 = df.copy()
from sklearn import preprocessing
le = preprocessing.LabelEncoder()
for column in df1.columns[1:]:
    df1[column]= le.fit_transform(df1[column])
df1.head()
```

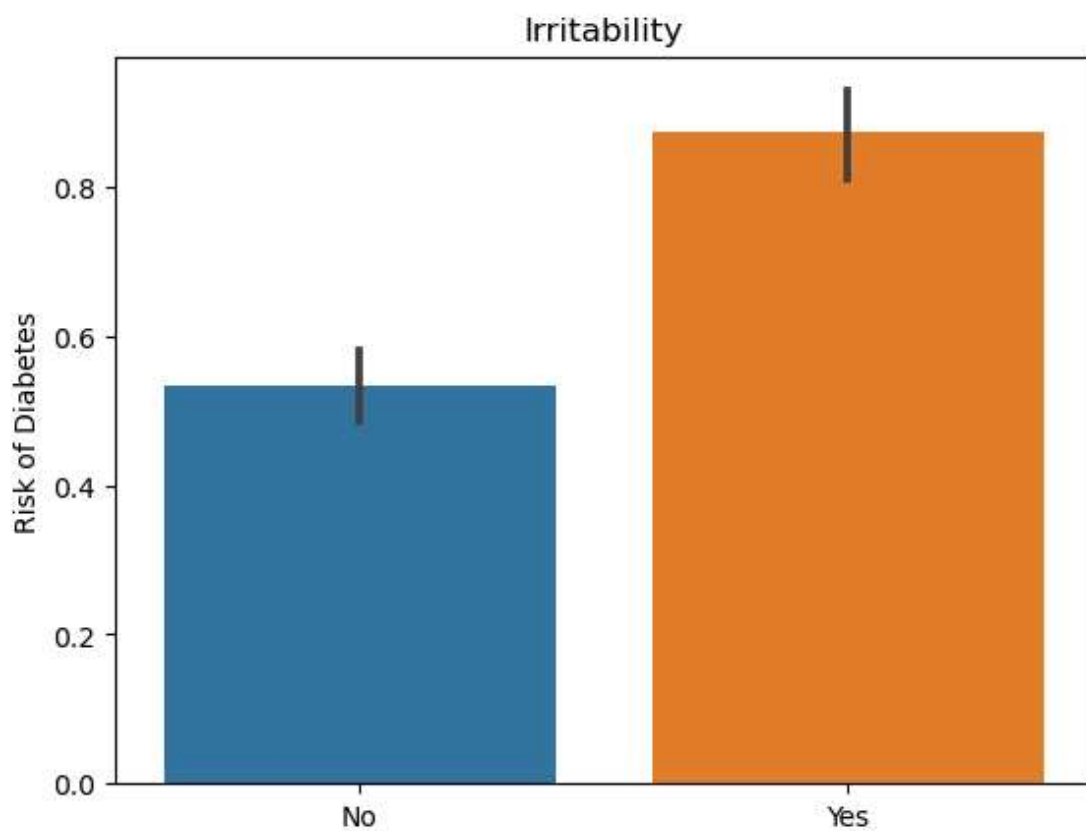
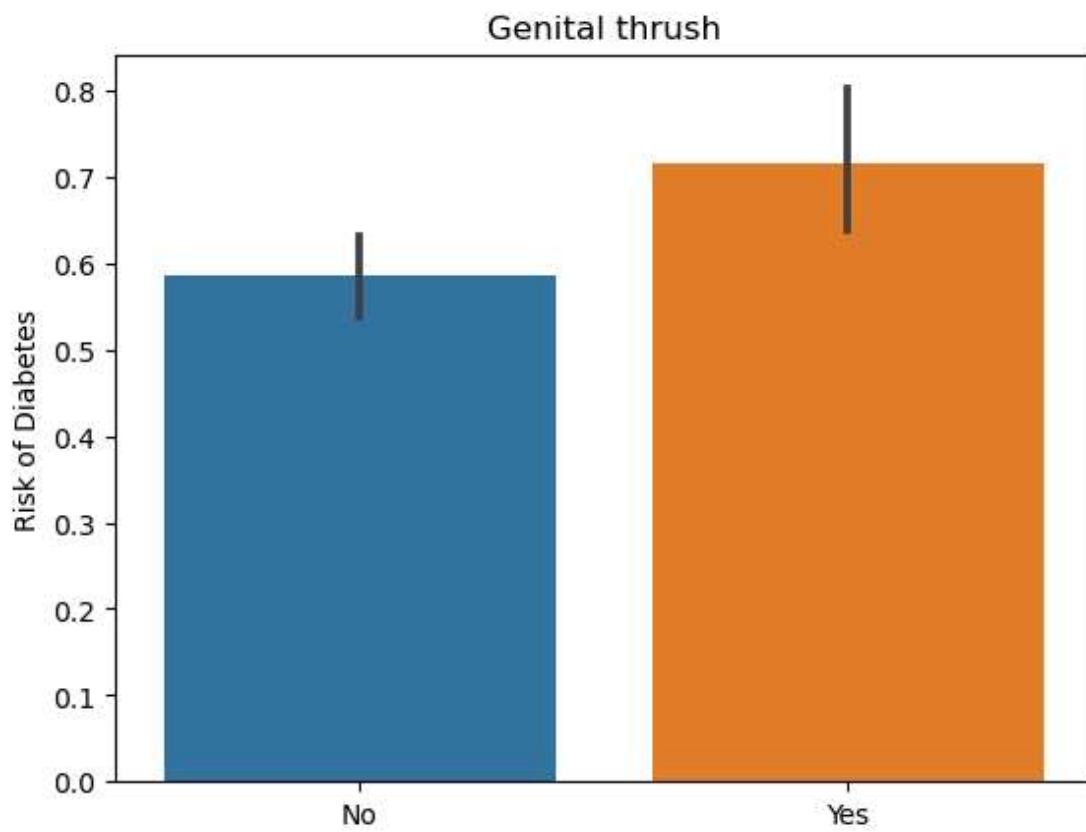
Out[6]:


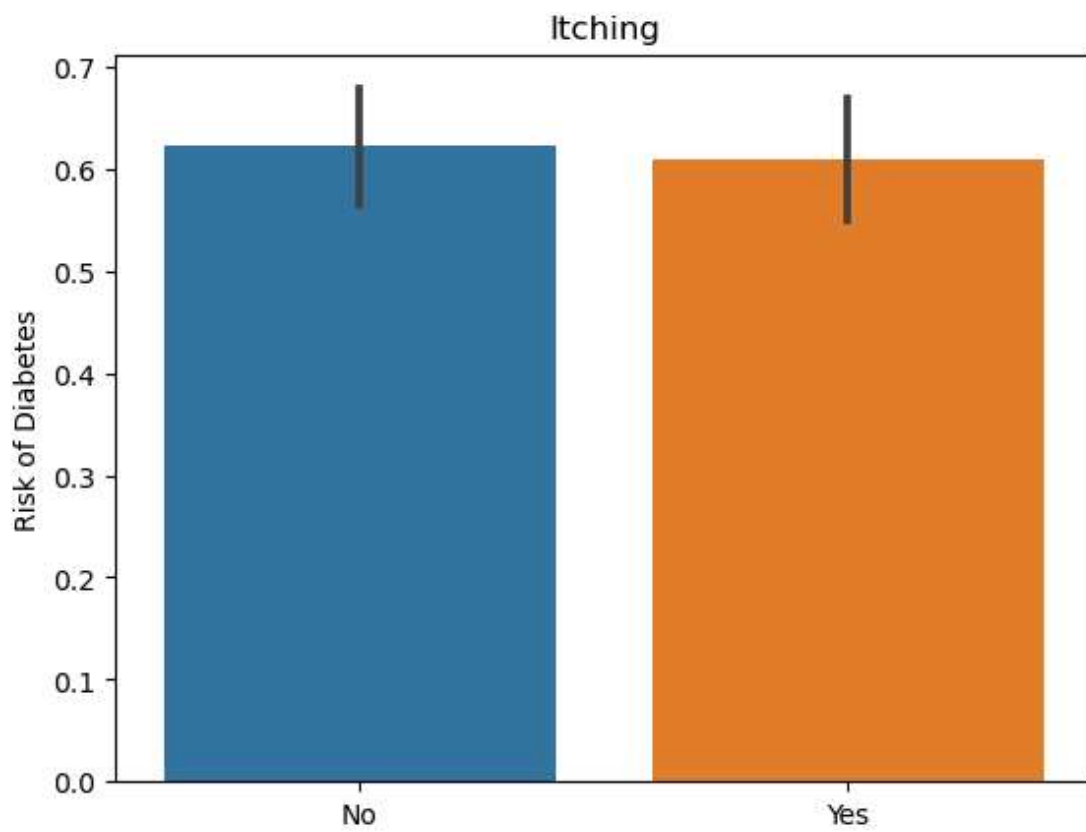
	Age	Gender	Polyuria	Polydipsia	sudden weight loss	weakness	Polyphagia	Genital thrush	visual blurring	Itching	Irritabili
0	40	1	0	1	0	1	0	0	0	1	
1	58	1	0	0	0	1	0	0	1	0	
2	41	1	1	0	0	1	1	0	0	1	
3	45	1	0	0	1	1	1	1	0	1	
4	60	1	1	1	1	1	1	0	1	1	


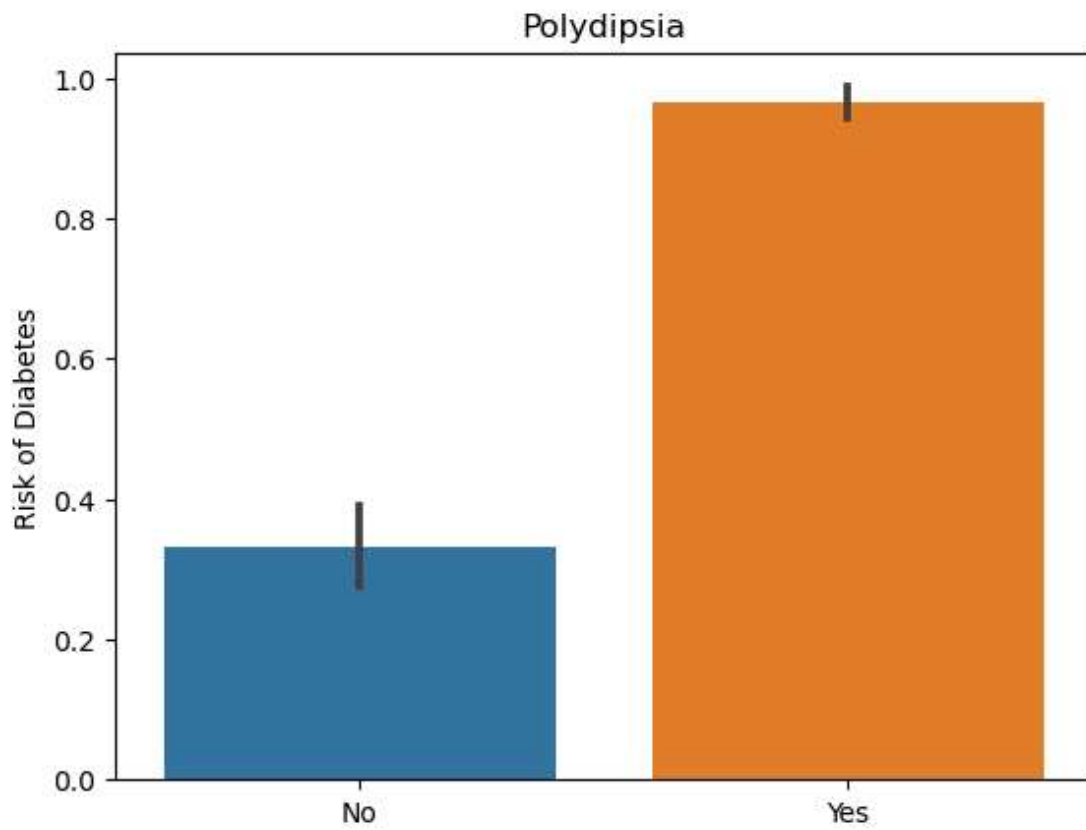
```
In [7]: import seaborn as sns
import matplotlib.pyplot as plt
df_symptoms = df1[df1.columns.difference(["Age", "class", "Gender"])]

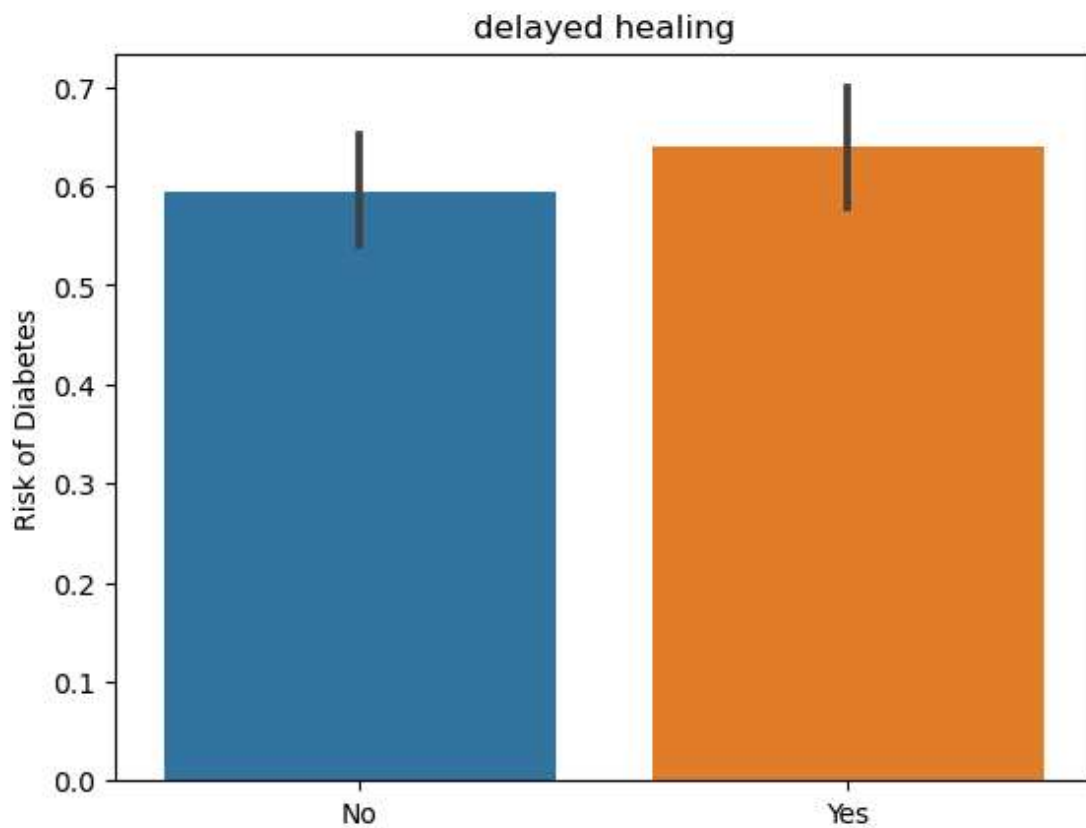
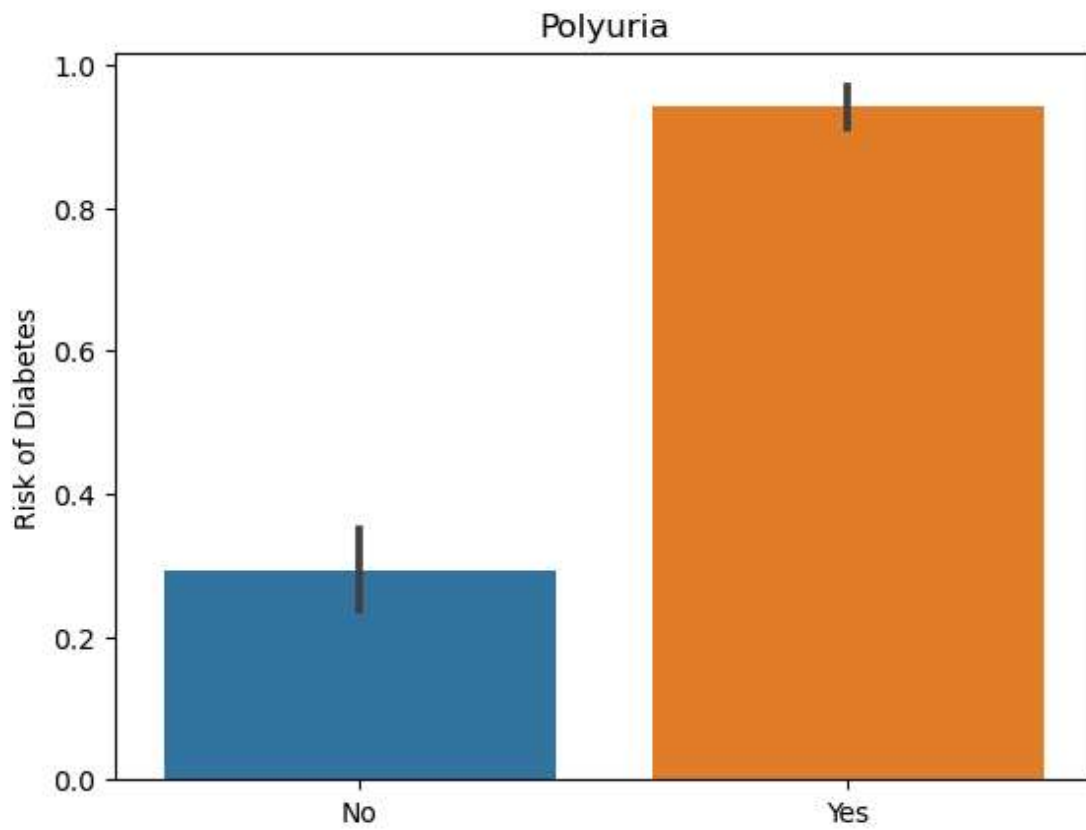
for column in df_symptoms.columns:
    ax = sns.barplot(x=column, y="class", data=df1)
    ax.set_xticklabels(["No", "Yes"])
    ax.set_ylabel("Risk of Diabetes")
    ax.set_xlabel(None)
    plt.title(column)
    plt.show()
```

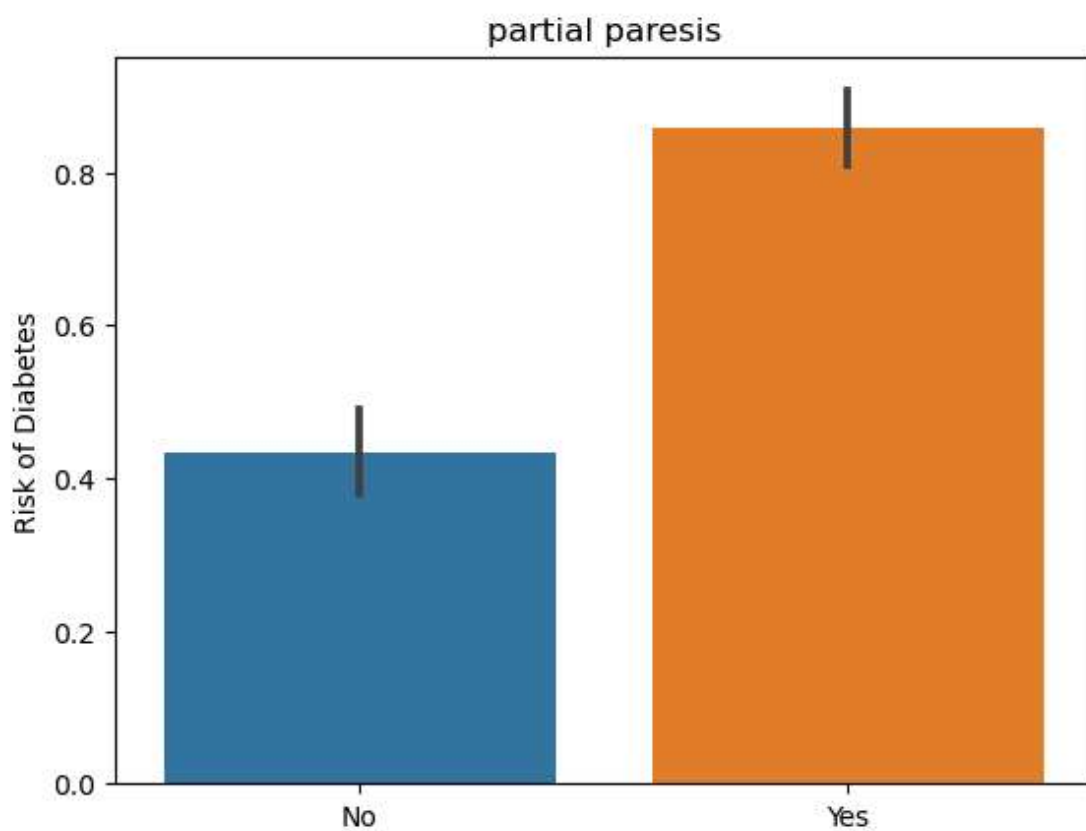
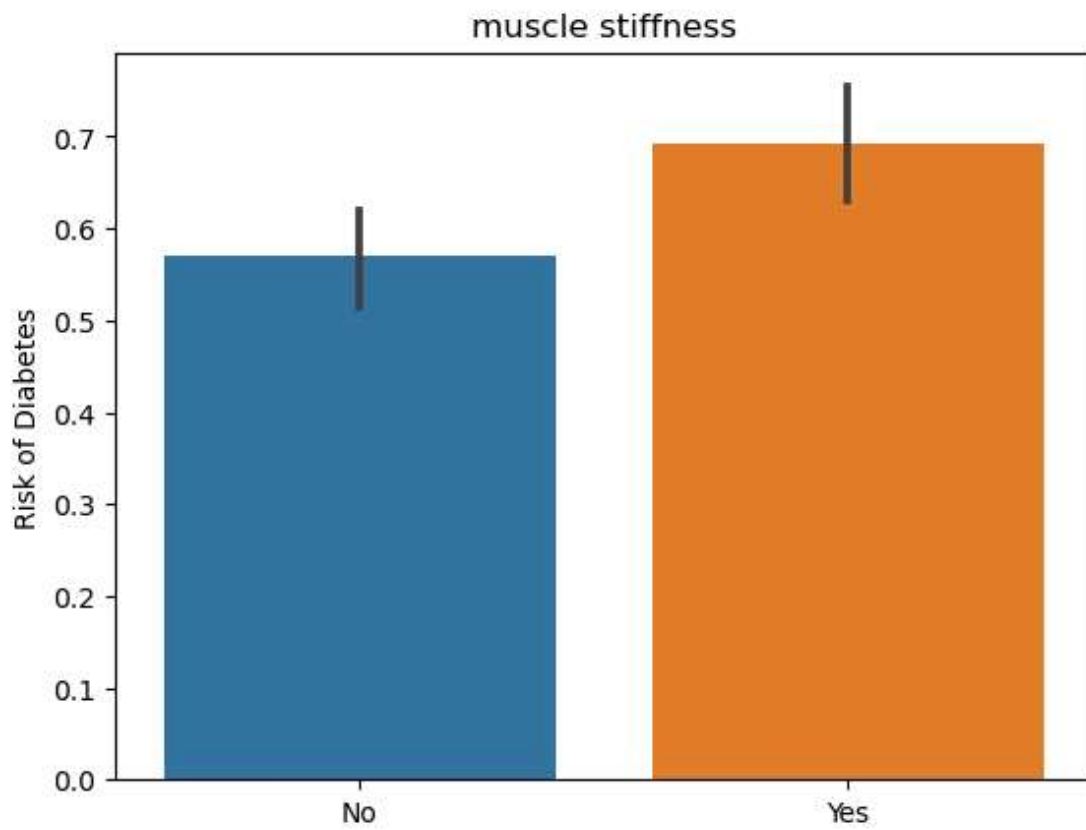


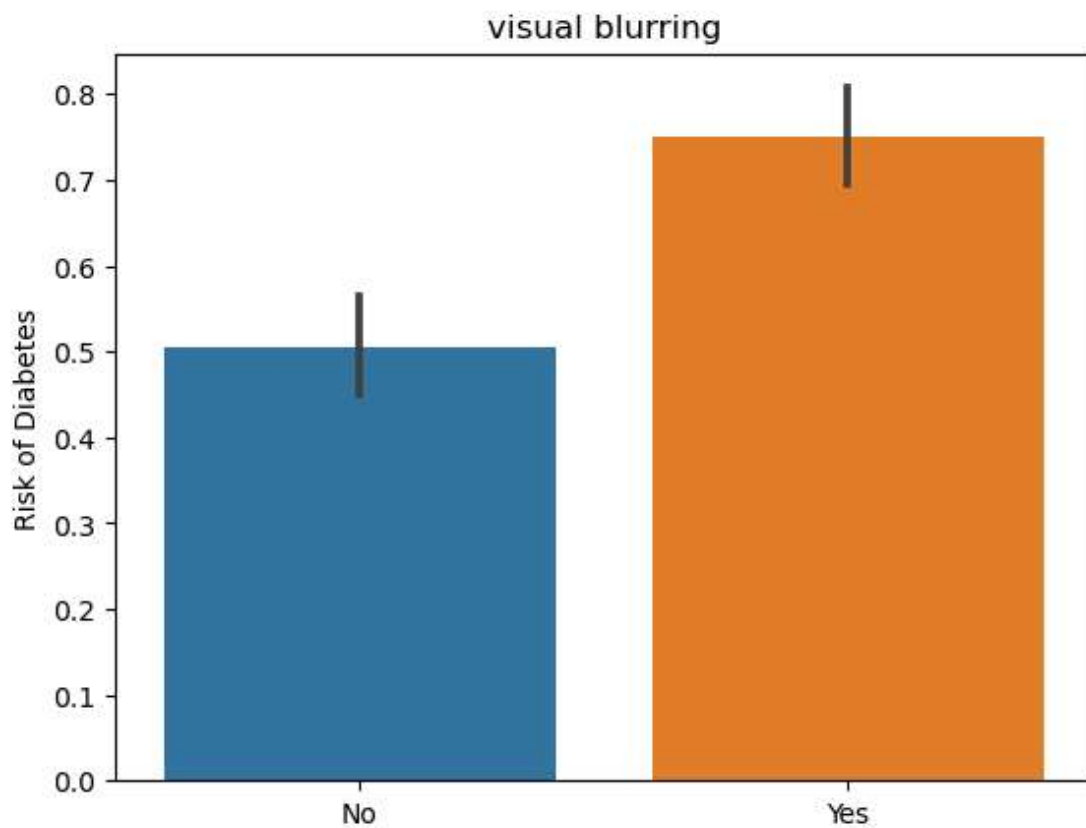
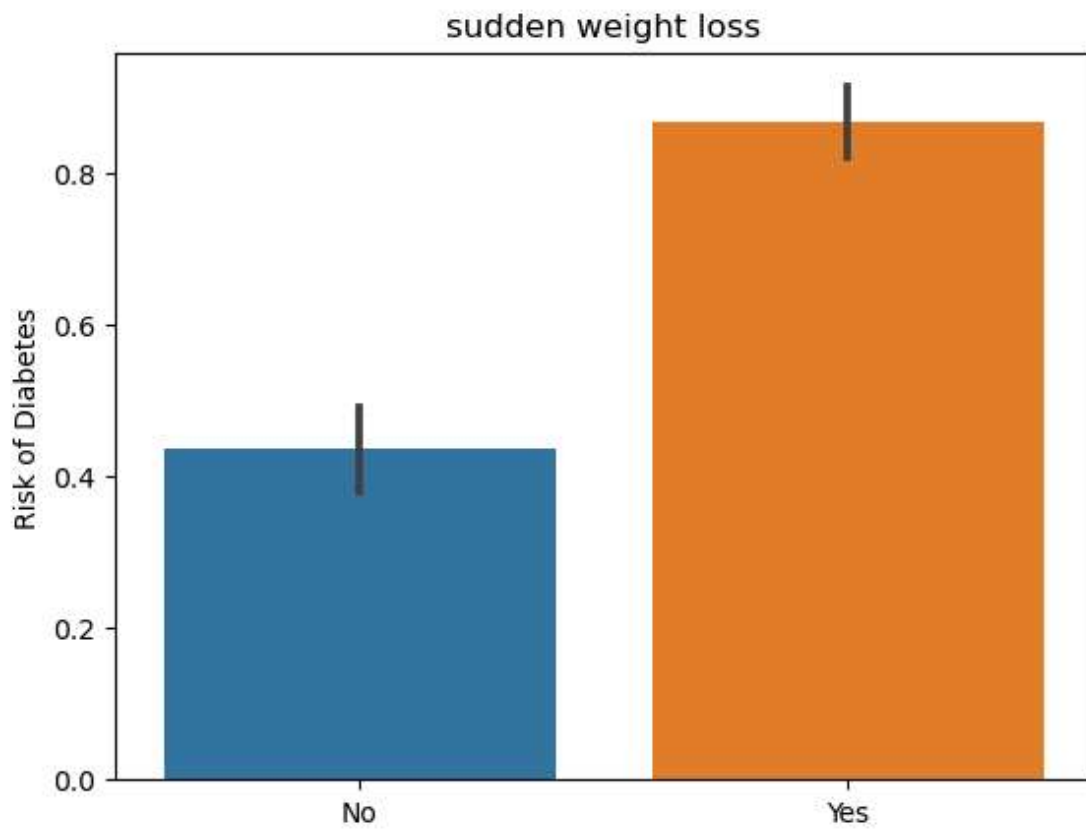

















```
In [8]: ag = sns.displot(df["Age"], color="c", kde= True)
```



```
In [9]: sns.heatmap(df1.corr(), cmap='RdYlGn')
```

```
Out[9]: <Axes: >
```



```
In [10]: subset =df.iloc[:, :-1]
sorted_vals = np.abs(subset.corrwith(df1['class'])).sort_values(ascending=False)
print(sorted_vals)
```

```
Age    0.108679
dtype: float64
```

C:\Users\shubh\AppData\Local\Temp\ipykernel_15732\1250707825.py:2: FutureWarning: The default value of numeric_only in DataFrame.corrwith is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
sorted_vals = np.abs(subset.corrwith(df1['class'])).sort_values(ascending=False)
```

```
In [11]: #copying the dataframe
X = df1.drop('class', axis =1)
y = df1['class']

#splitting the datasets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y , test_size = 0.3)
print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)
```

```
(364, 16) (364,)
(156, 16) (156,)
```

```
In [12]: from sklearn.model_selection import cross_val_predict, cross_val_score
from sklearn.model_selection import KFold, ShuffleSplit
from sklearn.preprocessing import StandardScaler, MinMaxScaler


#standard scaling
s=StandardScaler()
X_train_s = s.fit_transform(X_train)
X_test_s = s.transform(X_test)
```

```
In [13]: #Testing different models:
#Logistic Regression:
from sklearn.linear_model import LogisticRegression
LR = LogisticRegression()
LR = LR.fit(X_train_s,y_train)
y_predict = LR.predict(X_test_s)
print(f'model score on training data: {LR.score(X_train_s, y_train)}')
print(f'model score on testing data: {LR.score(X_test_s, y_test)}')
LR.coef_
```

```
model score on training data: 0.9505494505494505
model score on testing data: 0.8974358974358975
```

```
Out[13]: array([[ -0.18591246, -1.9070854 ,  1.77618485,  1.92788244,  0.26572276,
          0.70552922,  0.45347349,  0.3333031 ,  0.56431728, -1.61062041,
          1.0474483 ,  0.00828654,  0.48837376, -0.75434614,  0.0622248 ,
          -0.56989493]])
```

```
In [14]: from sklearn.metrics import confusion_matrix
sns.heatmap(confusion_matrix(y_test, y_predict), annot= True, fmt='d', cmap='Greens')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix for Logistic Regression')
plt.show()
```



```
In [15]: from sklearn.metrics import accuracy_score
from sklearn.inspection import permutation_importance

baseline_accuracy = accuracy_score(y_test, y_predict)

#permutation importance
perm_importance = permutation_importance(LR, X_test_s, y_test, n_repeats=10, random_state=42)

#feature importances
feature_importances = perm_importance.importances_mean
print("Permutation Importance:")
for feature_name, importance in zip(X.columns, feature_importances):
    print(f"{feature_name}: {importance:.4f} (Decrease in Accuracy)")
```

Permutation Importance:

Age: -0.0032 (Decrease in Accuracy)

Gender: 0.0808 (Decrease in Accuracy)

Polyuria: 0.1583 (Decrease in Accuracy)

Polydipsia: 0.1179 (Decrease in Accuracy)

sudden weight loss: -0.0045 (Decrease in Accuracy)

weakness: 0.0045 (Decrease in Accuracy)

Polyphagia: 0.0269 (Decrease in Accuracy)

Genital thrush: 0.0103 (Decrease in Accuracy)

visual blurring: 0.0128 (Decrease in Accuracy)

Itching: 0.0417 (Decrease in Accuracy)

Irritability: 0.0404 (Decrease in Accuracy)

delayed healing: 0.0000 (Decrease in Accuracy)

partial paresis: 0.0282 (Decrease in Accuracy)

muscle stiffness: 0.0019 (Decrease in Accuracy)

Alopecia: 0.0000 (Decrease in Accuracy)

Obesity: -0.0013 (Decrease in Accuracy)

```
In [16]: from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score
from sklearn.preprocessing import label_binarize
precision = precision_score(y_test, y_predict)
recall = recall_score(y_test, y_predict)
roc_auc = roc_auc_score(y_test, y_predict)
f1 = f1_score(y_test, y_predict)

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"ROC AUC: {roc_auc:.4f}")
print(f"F1 Score: {f1:.4f}")
```

Precision: 0.9318
Recall: 0.8913
ROC AUC: 0.8988
F1 Score: 0.9111

```
In [17]: #Linear SVC:
from sklearn.svm import LinearSVC
LinSVC = LinearSVC(C=10.0)
LinSVC = LinSVC.fit(X_train_s, y_train)
y_predict = LinSVC.predict(X_test_s)
print(f'model score on training data: {LinSVC.score(X_train_s, y_train)}')
print(f'model score on testing data: {LinSVC.score(X_test_s, y_test)}')
```


model score on training data: 0.9505494505494505
model score on testing data: 0.8974358974358975

C:\Users\shubh\anaconda3\Lib\site-packages\sklearn\svm_classes.py:31: FutureWarning: The default value of `dual` will change from `True` to `auto` in 1.5. Set the value of `dual` explicitly to suppress the warning.
warnings.warn(
C:\Users\shubh\anaconda3\Lib\site-packages\sklearn\svm_base.py:1237: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
warnings.warn(

```
In [18]: #SVC:
from sklearn.svm import SVC
rbfsvc = SVC(C=8, random_state=0)
rbfsvc = rbfsvc.fit(X_train_s, y_train)
y_predict = rbfsvc.predict(X_test_s)
print(f'model score on training data: {rbfsvc.score(X_train_s, y_train)}')
print(f'model score on testing data: {rbfsvc.score(X_test_s, y_test)}')
```

model score on training data: 0.9917582417582418
model score on testing data: 0.9487179487179487

```
In [19]: sns.heatmap(confusion_matrix(y_test, y_predict), annot=True, fmt='d', cmap='Greens')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()
```



```
In [20]: precision = precision_score(y_test, y_predict)
recall = recall_score(y_test, y_predict)
roc_auc = roc_auc_score(y_test, y_predict)
f1 = f1_score(y_test, y_predict)
```

```
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"ROC AUC: {roc_auc:.4f}")
print(f"F1 Score: {f1:.4f}")
```

```
Precision: 0.9884
Recall: 0.9239
ROC AUC: 0.9541
F1 Score: 0.9551
```

```
In [21]: #Random Forest:
from sklearn.ensemble import RandomForestClassifier
RC = RandomForestClassifier()
RC = RC.fit(X_train_s, y_train)
y_predict = RC.predict(X_test_s)
print(f'model score on training data: {RC.score(X_train_s, y_train)}')
print(f'model score on testing data: {RC.score(X_test_s, y_test)}')
```

```
model score on training data: 1.0
model score on testing data: 0.9615384615384616
```


```
In [22]: #permutation importance
perm_importance = permutation_importance(RC, X_test_s, y_test, n_repeats=10, random_state=42)

#feature importances
feature_importances = perm_importance.importances_mean
print("Permutation Importance:")
for feature_name, importance in zip(X.columns, feature_importances):
    print(f"{feature_name}: {importance:.4f} (Decrease in Accuracy)")
```


```
Permutation Importance:
Age: 0.0071 (Decrease in Accuracy)
Gender: 0.0609 (Decrease in Accuracy)
Polyuria: 0.1750 (Decrease in Accuracy)
Polydipsia: 0.1250 (Decrease in Accuracy)
sudden weight loss: 0.0058 (Decrease in Accuracy)
weakness: -0.0013 (Decrease in Accuracy)
Polyphagia: 0.0026 (Decrease in Accuracy)
Genital thrush: -0.0019 (Decrease in Accuracy)
visual blurring: -0.0026 (Decrease in Accuracy)
Itching: -0.0064 (Decrease in Accuracy)
Irritability: 0.0006 (Decrease in Accuracy)
delayed healing: 0.0006 (Decrease in Accuracy)
partial paresis: 0.0019 (Decrease in Accuracy)
muscle stiffness: -0.0071 (Decrease in Accuracy)
Alopecia: 0.0256 (Decrease in Accuracy)
Obesity: -0.0019 (Decrease in Accuracy)
```

```
In [23]: features = df1.drop('class', axis=1).columns
importances = RC.feature_importances_
indices = np.argsort(importances)

plt.figure()
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='r', align='center')
plt.yticks(range(len(indices)), [features[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```

```
In [24]: sns.heatmap(confusion_matrix(y_test, y_predict), annot= True, fmt='d', cmap='Greens')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix for Random Forest')
plt.show()
```



```
In [25]: precision = precision_score(y_test, y_predict)
recall = recall_score(y_test, y_predict)
roc_auc = roc_auc_score(y_test, y_predict)
f1 = f1_score(y_test, y_predict)
```

```
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"ROC AUC: {roc_auc:.4f}")
print(f"F1 Score: {f1:.4f}")
```

```
Precision: 0.9886
Recall: 0.9457
ROC AUC: 0.9650
F1 Score: 0.9667
```

```
In [26]: #Decision Trees:
from sklearn.tree import DecisionTreeClassifier
DTC = DecisionTreeClassifier(criterion = 'gini', max_features = 10 , max_depth = 5)
DTC = DTC.fit(X_train_s, y_train)
y_predict = DTC.predict(X_test)
print (f'model score on training data: {DTC.score(X_train_s, y_train)}')
print (f'model score on testing data: {DTC.score(X_test_s, y_test)}')

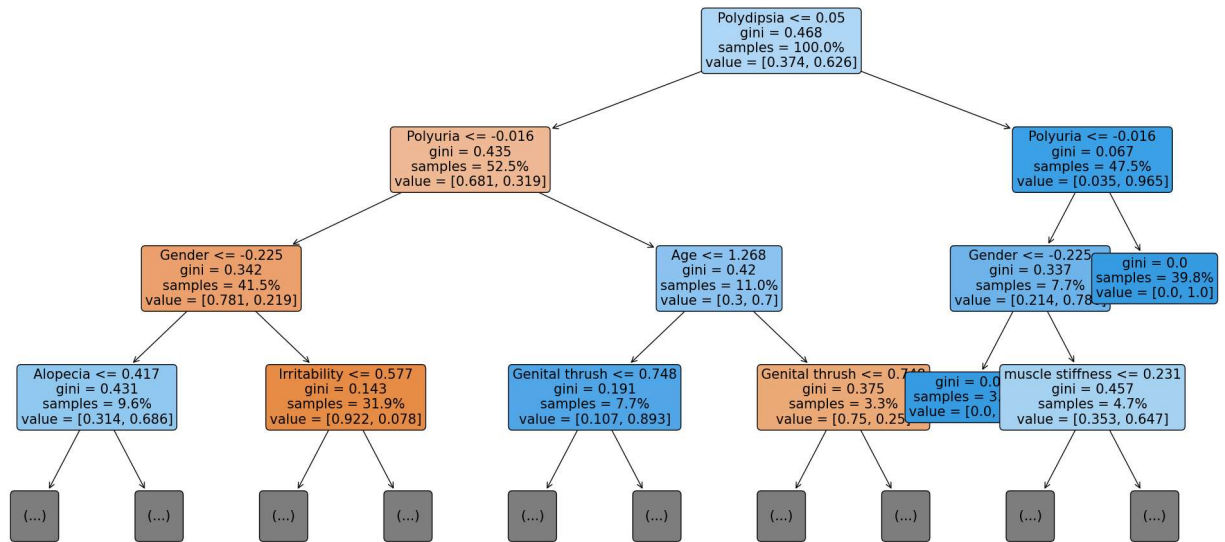
precision = precision_score(y_test, y_predict)
recall = recall_score(y_test, y_predict)
roc_auc = roc_auc_score(y_test, y_predict)
f1 = f1_score(y_test, y_predict)

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"ROC AUC: {roc_auc:.4f}")
print(f"F1 Score: {f1:.4f}")

model score on training data: 0.967032967032967
model score on testing data: 0.8846153846153846
Precision: 0.8182
Recall: 0.7826
ROC AUC: 0.7663
F1 Score: 0.8000
C:\Users\shubh\anaconda3\Lib\site-packages\sklearn\base.py:486: UserWarning: X has feature names, but DecisionTreeClassifier was fitted without feature names
  warnings.warn(
```

```
In [27]: from sklearn.tree import plot_tree
fig, ax = plt.subplots(figsize=(25,12))

plot_tree(
    decision_tree = DTC,
    feature_names = df1.columns,
    filled = True,
    rounded = True,
    proportion = True,
    max_depth = 3,
    fontsize = 15,
    ax=ax,
);
```



In []:

```

In [28]: #XGB Classifier
from xgboost import XGBClassifier
xgb = XGBClassifier()
xgb = xgb.fit(X_train_s, y_train)
y_predict = xgb.predict(X_test_s)
print(f'model score on training data: {xgb.score(X_train_s, y_train)}')
print(f'model score on testing data: {xgb.score(X_test_s, y_test)}')

precision = precision_score(y_test, y_predict)
recall = recall_score(y_test, y_predict)
roc_auc = roc_auc_score(y_test, y_predict)
f1 = f1_score(y_test, y_predict)

print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"ROC AUC: {roc_auc:.4f}")
print(f"F1 Score: {f1:.4f}")


model score on training data: 1.0
model score on testing data: 0.9551282051282052
Precision: 0.9885
Recall: 0.9348
ROC AUC: 0.9596
F1 Score: 0.9609

```

```

In [29]: sns.heatmap(confusion_matrix(y_test, y_predict), annot= True, fmt='d', cmap='Greens')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix for XGB Classifier')
plt.show()

```



```
In [30]: """
Different Models were tried for the dataset: Logistic Regression, Linear SVC, SVC, Random
Conclusion:
Random Forest works as the most accurate model which shows the best results with Polydipsia
(urine production by the body) being the most important features.
model score on training data: 1.0
model score on testing data: 0.9807692307692307
Precision score: 0.9895
Recall score: 0.9792
ROC AUC: 0.9812
F1 score: 0.9843
"""
```

```
Out[30]: '\nDifferent Models were tried for the dataset: Logistic Regression, Linear SVC, SVC, Ra
ndom Forest, Decision Trees, XGB Classifier\nConclusion:\nRandom Forest works as the mos
t accurate model which shows the best results with Polydipsia(excessive thirst) and Polyu
ria(excessive \nurine production by the body) being the most important features.\nmodel
score on training data: 1.0\nmodel score on testing data: 0.9807692307692307\nPrecision
score: 0.9895\nRecall score: 0.9792\nROC AUC: 0.9812\nF1 score: 0.9843\n'
```